

The messages list in your code snippet is used to manage the conversation history with the Azure OpenAI API in the `run_session` function

This is the initial state of the messages list, created at the start of the `run_session` function before any API calls are made. Specifically, it occurs in this line:

This happens immediately after the MCP client session is initialized (`await session.initialize()`) and before the first Azure OpenAI API call (`oai_client.chat.completions.create`).

```
messages = [
    {
        "role": "system",
        "content": "You have access to two tools via MCP: search_by_image and image_search..."
    },
    {
        "role": "user",
        "content": "Find 5 similar images for this image: https://cloudinary.com/image123.jpg"
    }
]
```

This represents the evolved state of the messages list after one or more iterations of processing user queries, tool calls, and model responses.

```
messages = [
    {"role": "system", "content": "prompt."},
    {"role": "user", "content": "Find images like this: url1"},
    {"role": "assistant", "tool_calls": [...]},
    {"role": "tool", "content": "{results}"},
    {"role": "assistant", "content": "Here are similar images..."},
    {"role": "user", "content": "Now find images like this: url2"}, # New request
    # AI remembers previous context
]
```

the `msg` object, which is a `ChatCompletionMessage` extracted from the Azure OpenAI API response in the `run_session` function. This specific output occurs when the model decides to invoke tools (`search_by_image` and `image_search`) in response to the user's query. Let's break down when and under what conditions this exact `msg` output is generated.

```
print(msg)
ChatCompletionMessage(content=None, role='assistant', tool_calls=[ToolCall(id='call_abc123',
type='function', function=FunctionCall(name='search_by_image',
arguments={'query_image_path': "https://example.com/image.jpg", "k": 5})),
ToolCall(id='call_def456', type='function', function=FunctionCall(name='image_search',
arguments={'image_url': "https://example.com/image.jpg", "num": 5})))])
```

Why mcp?

- 1) If one tool fails, we can still run the rest of the tools
- 2) If we want to migrate from azure open ai to some other service, its easy because server would stay the same
- 3) Reusable tools
- 4) Llm smarty decides if or not to call the tools
- 5) Can easily add new tools because ecosystem is not fragmented
- 6) The code uses asyncio and MCP's ClientSession to handle tool calls asynchronously (await session.call_tool(name, args)). This is critical for tools like image_search, which may involve slow network requests to external APIs (e.g., SerpAPI).

After tool execution, tool_payloads might look like:

```
tool_payloads = {
    "search_by_image": [
        {"decoded_path": "/local/db/image1.jpg", "description": "Red apple", "score": 0.92},
        {"decoded_path": "/local/db/image2.png", "description": "Fruit basket", "score": 0.85}
    ],
    "image_search": [
        {"title": "Red Apple Stock Photo", "url": "https://example.com/red-apple.jpg"},
        {"title": "Fruit Gallery", "url": "https://stock.example.com/fruit.png"}
    ]
}
```